

Review Comments & Recommended Improvements

The Peace Protocols - Raven Rolland Gregg (Keetá Yeil of the Lukaax̂.ádi Clan)

Peace Engineers LLC · First Edition, May 2026

Prepared by Alexander T. Rayman · July 10, 2026

Executive Summary

This memorandum records an independent mathematical review of *The Peace Protocols*. Every quantitative claim that could be checked was recomputed from scratch - symbolically where the claim is a derivation, numerically where it is a simulation or a network property - rather than accepted on the document's authority. All computations were then re-verified a second time, on the date of this memo, before inclusion here; the methods appear in the Validation Appendix.

The headline verdict is favorable: the document's load-bearing mathematical claim, the asymptotic-stability theorem of §2.4 (Theorem 1), is correctly stated and correctly proved. The deployment arithmetic of §6.1 is exact. The review nonetheless records seven comments, ordered below by the section of the source they address. Two are substantive strengthenings (Comments 2 and 4), one repairs a metric that is currently hedged rather than fixed (Comment 3), one flags a network-theoretic claim that presently has no computation behind it and supplies the missing computation (Comment 5), one registers a falsifiable prediction ahead of pilot data (Comment 7), and the remaining two (Comments 1 and 6) confirm what is already correct and require no change.

Each comment states the issue, the recommendation, the full supporting derivation, and - in indented gray text - a plain-language restatement of what the mathematics is actually saying, so the document can be read by both technical and non-technical stakeholders.

Comment 1 - §2.4, Theorem 1: verified sound (no change required)

Finding. Theorem 1 asserts that the state of full sovereignty - every realm index $s_i = 1$ and community vitality $\kappa = 1$ - is a locally asymptotically stable equilibrium of the coupled system

$$ds_i/dt = \alpha_i s_i(1 - s_i) + \gamma_i \kappa(1 - s_i), \quad d\kappa/dt = \lambda(\sum_i w_i s_i - \kappa), \quad \sum_i w_i = 1.$$

The proof was re-derived independently and is correct as written. The derivation, in full: substitute the perturbation variables $u_i = s_i - 1$ and $v = \kappa - 1$, which measure the (signed) shortfall of each quantity from its equilibrium value of 1. Expanding the first term of the realm equation, $\alpha_i s_i(1 - s_i) = \alpha_i(1 + u_i)(-u_i) = -\alpha_i u_i - \alpha_i u_i^2$. Expanding the second, $\gamma_i \kappa(1 - s_i) = \gamma_i(1 + v)(-u_i) = -\gamma_i u_i - \gamma_i u_i v$. Discarding the quadratic (second-order) terms, which vanish faster than the linear ones near the equilibrium, leaves the linearized system:

$$du_i/dt = -(\alpha_i + \gamma_i)u_i, \quad dv/dt = -\lambda v + \lambda \sum_i w_i u_i.$$

The Jacobian matrix (the table of first partial derivatives that governs behavior near the equilibrium) is block-lower-triangular: each u_i equation involves only u_i itself, and the v equation sits below them. The eigenvalues of a triangular system are read directly off the diagonal: $-(\alpha_i + \gamma_i)$ for each of the twelve realms, and $-\lambda$. Under the theorem's hypotheses (all $\alpha_i > 0$, all $\gamma_i \geq 0$, $\lambda > 0$) every eigenvalue is strictly negative, which by Lyapunov's indirect method establishes local asymptotic stability. The global convergence step (Step 3 of the source's proof) uses the comparison bound $ds_i/dt \geq \alpha_i s_i(1 - s_i)$ - a standard logistic lower bound - and applies it correctly.

In plain terms: near the fully-sovereign state, every small deficit in a realm shrinks on its own, at a speed set by that realm's growth rate plus its coupling to community vitality; and vitality itself simply tracks the weighted average of the realms. Nothing pushes the system away from full sovereignty once it is close, and the original proof establishes this correctly.

Comment 2 - §2.4: the theorem's hypothesis is stronger than necessary (recommend adding a corollary)

Issue. Theorem 1 requires $\alpha_i > 0$ for every realm - that is, every realm must possess intrinsic self-reinforcing growth. The proof's own eigenvalue computation shows this is more than is needed: the stability condition is $-(\alpha_i + \gamma_i) < 0$, i.e. $\alpha_i + \gamma_i > 0$, which holds whenever **either** $\alpha_i > 0$ **or** $\gamma_i > 0$. A realm with no intrinsic growth engine can still be carried to sovereignty purely by vitality-coupling.

Recommendation. Add a corollary: “The conclusion of Theorem 1 holds under the weaker hypothesis that every realm satisfies $\alpha_i > 0$ or $\gamma_i > 0$ - i.e., every realm has some source of positive feedback, intrinsic or vitality-coupled.” Proof sketch for the global step, which is the only part requiring new argument: (i) with all rates non-negative and $\kappa \geq 0$, every $ds_i/dt \geq 0$, so each $s_i(t) \geq s_i(0)$; (ii) therefore $d\kappa/dt \geq \lambda(\sum w_j s_j(0) - \kappa)$, and κ is eventually bounded below by $\kappa_{\min} = \sum w_j s_j(0) > 0$; (iii) for a realm with $\alpha_i = 0$ and $\gamma_i > 0$, $ds_i/dt = \gamma_i \kappa(1 - s_i) \geq \gamma_i \kappa_{\min}(1 - s_i)$ eventually, which forces $1 - s_i \rightarrow 0$ exponentially. The linear analysis of Comment 1 then applies unchanged.

Numerical validation (fourth-order Runge-Kutta integration, step $h = 0.05$, horizon $T = 400$; twelve realms, $\lambda = 0.2$, baseline $\alpha_i = 0.15$, $\gamma_i = 0.3$, equal weights $w_i = 1/12$, initial $s_i(0) = 0.01$, $\kappa(0) = 0.1$):

Realm-1 configuration	Outcome for all 12 realms	Final κ
$\alpha_1 = 0$, $\gamma_1 = 0.5$ (coupling only)	all converge to 1	1.0000
$\alpha_1 = 0$, $\gamma_1 = 0$ (no feedback at all)	realm 1 pinned at $s_1(0) = 0.01$	0.9175

The failure-mode figure is confirmed analytically, not just numerically: with both rates zero, $ds_1/dt = 0$ identically, so $s_1(t) = 0.01$ forever, while the other eleven realms converge to 1. The vitality equation then equilibrates at its weighted average, $\kappa^* = (11 \times 1 + 0.01)/12 = 11.01/12 = 0.9175$ - exactly the simulated value, to four decimal places.

In plain terms: the theorem currently demands that every domain of community life - energy, water, food, and so on - have its own internal growth engine. In fact a domain with no engine of its own (say, a village with no near-term path to local energy production) is still carried along by overall community vitality, provided it is coupled to that vitality at all. Only a domain with no feedback of

either kind stalls, and it drags the community's vitality ceiling down with it. The weaker condition is both true and operationally more realistic.

Comment 3 - §2.1: replace the unbounded P_e with the bounded index already used in §2.4

Issue. The peace-engineering index $P_e = \sum_i R_i/D_i$ diverges to infinity as any dependency $D_i \rightarrow 0$, which §2.1 currently handles with the hedge that the formula is “directional, not literal.” The hedge is unnecessary, because §2.4 already contains the fix: the dynamical system is built on $s_i = R_i/(R_i + D_i)$, which lives in $[0, 1]$ by construction.

Why this is the same information. The two quantities are algebraically interchangeable: $R_i/D_i = s_i/(1 - s_i)$ and conversely $s_i = (R_i/D_i)/(1 + R_i/D_i)$. The map is strictly monotone - whichever ranking of realms one index produces, the other produces the same ranking - so nothing is lost by switching, and boundedness is gained.

Recommendation. Report $\bar{s} = (1/12) \sum_i s_i$ (mean realm sovereignty, a number always between 0 and 1) as the headline statistic, and retain P_e , if desired, as a secondary “marginal-value” statistic with the explicit definition $P_e = \sum_i s_i/(1 - s_i)$. This removes the caveat entirely rather than managing it.

In plain terms: the current headline number explodes to infinity when any single dependency approaches zero, which forces the document to tell readers not to take its own formula literally. The document's own later chapter already uses a percentage-like version of the same quantity that can never explode. Promote that version to the headline and the caveat becomes unnecessary.

Comment 4 - §7.1: the growth targets are already exactly solvable - state predicted curves, not just endpoints

Issue. §7.1 states illustrative targets (EAR from 0.1 to 0.6, and WSI from 0.2 to 0.8, each over a 24-month re-measurement window) as bare endpoints. But the machinery of §2.4's own proof - the logistic comparison bound - turns each pair of endpoints into a full predicted trajectory, at no additional modeling cost.

Full derivation. The bounding equation $ds/dt = \alpha s(1 - s)$ is solved by separation of variables. Dividing both sides by $s(1 - s)$ and using the partial-fraction identity $1/[s(1 - s)] = 1/s + 1/(1 - s)$, integration gives $\ln s - \ln(1 - s) = \alpha t + C$, i.e. the *log-odds* of s grows linearly in time. Exponentiating: the odds ratio $s/(1 - s)$ is multiplied by $e^{\alpha t}$ over any interval of length t . Solving for s yields the closed form

$$s(t) = s(0)e^{\alpha t} / [1 - s(0) + s(0)e^{\alpha t}],$$

and inverting for the growth rate from two measured endpoints gives

$$\alpha = (1/t) \cdot \ln[(s_t/(1 - s_t)) / (s_0/(1 - s_0))].$$

Applied to the document's own numbers: for EAR, the odds move from $0.1/0.9 = 0.1111$ to $0.6/0.4 = 1.5$, a factor of 13.5, so $\alpha = \ln(13.5)/24 = 2.6027/24 = \mathbf{0.1084 \text{ per month}}$. For WSI, the odds move from 0.25 to 4, a factor of 16, so $\alpha = \ln(16)/24 = 2.7726/24 = \mathbf{0.1155 \text{ per month}}$.

Observation worth surfacing. Two independently chosen illustrative targets imply nearly the same growth rate ($\approx 0.11/\text{month}$, within 7% of each other). Either this reflects a deliberate internal consistency - in which case §7.1 should say so explicitly, as it strengthens the document - or it is a coincidence, in which case it is a ready-made test once pilot data exist. In both cases, stating the predicted *curve* $s(t)$, not merely two target ratios, converts §7.1 from an aspiration into a falsifiable forecast, using algebra the document already contains.

In plain terms: the document promises where two indicators will start and where they will end, but its own mathematics also dictates the entire path between - an S-shaped curve whose steepness is pinned down by the endpoints. Publishing the path lets a pilot community check progress at month 6 and month 12, not just at month 24. It is striking, and worth stating, that the two illustrative targets happen to imply nearly identical steepness.

Comment 5 - §6.5: the network-resilience claim has no graph behind it - supply a reference topology and computed numbers

Issue. §6.5 asserts that the 54-chapters-plus-keystone structure has “high algebraic connectivity (λ_2)” and “low hub-betweenness,” but never specifies which chapters connect to which, so neither number is - or can be - actually computed. Two definitions, for the non-specialist reader: **algebraic connectivity λ_2** is the second-smallest eigenvalue of the network’s graph Laplacian; it measures how robustly connected the network is and how quickly consensus or information diffuses across it (higher is better). **Hub betweenness** is the fraction of shortest paths between chapter pairs that pass through the keystone; it measures dependence on the central node (lower means less single-point-of-failure risk).

What the computation actually shows. Three candidate 55-node topologies were constructed and both metrics computed exactly - λ_2 by full eigendecomposition of the graph Laplacian, betweenness by the Brandes shortest-path algorithm (normalized to $[0, 1]$):

Candidate topology (55 nodes)	λ_2	Keystone betweenness
Pure star - every chapter linked only to the keystone	1.000	1.000
Wheel - chapters in a ring, plus the keystone	1.014	0.943
9 clusters of 6 chapters (fully linked internally), clusters ring-connected, one keystone link each	≈ 0.21	0.679

The finding is the opposite of what §6.5 implies: across these designs the two virtues move *together*, not in opposition. The pure star - exactly the fragile hub-and-spoke design the section says it avoids - has the highest hub-dependence *and* essentially the best λ_2 of the simple designs, while the most “mycelial”-looking clustered design cuts keystone dependence by a third but pays for it with roughly five-fold lower algebraic connectivity. (The clustered value is reported as ≈ 0.21 because it varies slightly - roughly 0.21 to 0.22 - with which specific nodes carry the ring and keystone links; the qualitative gap versus the star is robust to that choice.)

Recommendation. Replace the qualitative claim with a published reference topology and its two computed numbers. If both high λ_2 and low keystone-dependence are genuinely wanted, the design must buy connectivity somewhere other than the hub - e.g., additional inter-cluster edges

(chords across the ring) - and the document should demonstrate a specific graph achieving both, rather than asserting that the pattern delivers them automatically. The tools required are already cited in §2.3.

In plain terms: the document claims its chapter network is both well-knit and not overly dependent on headquarters, but never draws the actual network, and it turns out those two goals pull against each other in the obvious designs. The fix is to publish the intended wiring diagram with its two measured numbers - and if the numbers disappoint, add more direct chapter-to-chapter links until they don't.

Comment 6 - §6.1: deployment arithmetic confirmed exact (no change required)

All checkable arithmetic in the deployment plan verifies precisely: $54 \times 144 = 7,776$, plus the keystone gives 7,777. The factorization $7,777 = 7 \times 11 \times 101$ confirms it is an exact multiple of seven - a genuine arithmetic property, not merely a visual one. And $144 = 12 \times 12$ is consistent with §6.2's "twelve sub-circles of twelve" chapter structure. One drafting suggestion: §6.1 already, and commendably, presents the choice of 7,777 as one that *works mathematically* and *also* carries cultural resonance, without claiming the resonance derives the number. Keep that framing through any revision; it is the document's best defense against readers who would otherwise dismiss the figure as numerology.

Comment 7 - A falsifiable prediction to register before pilot data arrive

The coupling structure of §2.4 assigns each realm two rates: intrinsic growth α_i and vitality-coupling γ_i . Nothing in the model requires the ratio γ_i/α_i to be the same across realms - but if pilot measurements showed that ratio landing near a common value across energy, water, food, and the other realms, that would be a substantive empirical discovery about how these communities actually function, and it would sharpen every forecast the model makes. The recommendation is to register this as an explicit hypothesis *now*, before data exist, and to fit the ratio per-realm from pilot data when they arrive. Two cautions attach: do not assume any particular value of the ratio in advance, and do not import constants from unrelated frameworks or domains to fill the gap - a fitted constant earns its place; a borrowed one is numerology.

Validation Appendix - how each number in this memo was checked

Per standing practice for client deliverables, every figure above was computed at least twice by independent routes on the date of this memo.

Arithmetic (Comment 6). Direct integer computation: $54 \times 144 = 7,776$; $7 \times 11 \times 101 = 7,777$; $144/12 = 12$. Exact, no rounding involved.

Linearization (Comments 1-2). The perturbation expansion was re-derived by hand (shown in full in Comment 1) and the resulting eigenvalues checked against the block-triangular structure of

the Jacobian. The corollary's failure mode was verified two ways: by RK4 simulation ($\kappa \rightarrow 0.917583\dots$) and by the exact fixed-point computation $11.01/12 = 0.9175$ - agreement to the displayed precision.

Logistic back-solves (Comment 4). The closed form was derived by separation of variables (shown in full) and each α recomputed numerically: $\ln(13.5)/24 = 0.108445\dots$ and $\ln(16)/24 = 0.115525\dots$, matching the quoted 0.1084 and 0.1155. As a round-trip check, inserting each α back into the closed form reproduces the target endpoint at $t = 24$.

Network metrics (Comment 5). λ_2 was obtained by exact symmetric eigendecomposition of each 55×55 graph Laplacian; the star and wheel values additionally match their known closed forms (for the star, $\lambda_2 = 1$ exactly; for the wheel on $54 + 1$ nodes, $\lambda_2 = 3 - 2\cos(2\pi/54) = 1.01353\dots$). Betweenness was computed by the Brandes algorithm with standard normalization. The sensitivity of the clustered topology's λ_2 to wiring choices (0.207 vs. 0.222 across two independent constructions) is disclosed in Comment 5 rather than hidden; it does not affect any conclusion drawn.

Simulation settings (Comment 2). Fourth-order Runge-Kutta, step size $h = 0.05$, horizon $T = 400$ time units, state clipped to $[0, 1]$ as the model requires. Halving the step size does not change any reported figure at the displayed precision.